

Evan Reilly

✉ evanreillyswe@gmail.com 📞 (608) 728 1207
🌐 evanreilly.me 🐙 [/ereilly89](https://github.com/ereilly89) 🌐 [/evan-reilly](https://www.linkedin.com/in/evan-reilly)

Software Engineer

Software engineer with 4+ years of experience in distributed systems, Kubernetes-based platforms, and cloud-native architectures. My strengths include designing scalable backend services, building reliable cloud-native infrastructure, and developing distributed systems that support modern AI-enabled applications.

Technical Skills

Languages: Python, Java, Bash, JavaScript, SQL, Go
Technologies: Kubernetes, OpenShift, Git, Docker, Node.js, React, MySQL, MongoDB, Django, Scikit-learn, TensorFlow, Apache Kafka
Infrastructure: IBM Cloud (IKS, ROKS, COS), AWS (EC2, ELB, EBS, S3, Route53, CloudWatch), GCP (GAE, Vertex AI), Ansible, Jenkins, Linux, IAM
AI Systems: RAG agents, MCP, LlamaIndex, NLP pipelines
Developer Tools: Claude, Cursor, ChatGPT, Microsoft Copilot

Experience

IBM Cloud Software Engineer II Jan 2022 – Dec 2025

- Architected backend services for a 3D BI platform, earning recognition from IBM’s **Academy of Technology**
- Automated CPU and memory profiling for Kubernetes control plane components, accelerating resource tuning by **10×**
- Engineered CI/CD automation for Kubernetes and OpenShift services, reducing validation effort by **20+ dev hours monthly**
- Built an NLP + TF-IDF pipeline to cluster Kubernetes support tickets, **doubling triage efficiency**
- Extended Prometheus and Grafana observability tooling, reducing manual operational workload by **24+ hours monthly**
- Automated control plane warning and error detection, accelerating component verification by **5×**
- Resolved **150+** complex Kubernetes platform support tickets through deep technical debugging and customer collaboration
- Triageed and debugged **300+** Kubernetes and OpenShift end-to-end release test failures over 12 months
- Reviewed **300+** pull requests across Kubernetes and OpenShift service codebases, ensuring reliability and release readiness
- Built and maintained **24+** Kubernetes and OpenShift service release pipelines supporting production deployments
- Awarded **Outstanding Technical Achievement** for enabling migration of OpenShift clusters to HyperShift architecture

IBM Cloud Software Engineer Intern May 2021 – Aug 2021

- Automated cluster diagnostics with Go, Jenkins, and Slack bot integration, reducing ticket response time by **56+ hours per month**

Acuity Insurance Software Engineer Intern May 2020 – Aug 2020

- Built a full-stack catalog platform using JavaScript (Dojo), Java, and PostgreSQL enabling analysis of **300K+ insurance policies**
- Modernized legacy mainframe workflows to an Apache Kafka Pub/Sub architecture, increasing underwriting productivity by **25%**

U.S. Naval Research Laboratory ML Research Intern May 2019 – Aug 2019

- Developed a StarCraft II reinforcement learning agent in TensorFlow, achieving **15x faster** convergence than baseline models
- Designed MySQL relational database powering a leaderboard application for tracking reinforcement learning agent performance

UW Whitewater Research Assistant Sep 2018 – Dec 2021

- Published an explainable AI method for reinforcement learning, producing sequences **364% more informative** than baseline
- Developed a Java, JSP, and MySQL web application visualizing career statistics for **300+ computer science alumni**
- Delivered technical sessions on database systems and AI to **100+ students** improving practical engineering understanding
- Analyzed **50 Naval maintenance SOP trajectories** using Inverse Reinforcement Learning, optimizing time and cost

Projects

Fit to Scale Node.js, Express, React, MySQL, MongoDB, Docker, AWS, EC2, S3, Route53, CloudWatch, Stripe API

- Architected a multi-tenant SaaS platform enabling **10+ trainers** to sell subscription-based fitness services
- Built a microservices backend with APIs, SQL, and document databases processing **3.5M+ events annually**
- Built a React frontend for trainers to manage workouts, clients, subscriptions, and analytics dashboards

Kubernetes RAG Agent Python, LlamaIndex, ChromaDB, Claude, Streamlit, Docker

- Built a RAG pipeline ingesting **30 official Kubernetes docs**, enabling hallucination-free answers from verified source material
- **Eliminated embedding API costs** by running model locally, persisting vectors in ChromaDB for instant index reloads

Education

M.S. University of Wisconsin – Whitewater (Computer Science | 4.0 GPA)
• **Coursework:** Distributed Systems, Machine Learning, Advanced Database Systems, Advanced Algorithm Design & Analysis

B.S. University of Wisconsin – Whitewater (Computer Science, Economics | 3.97 GPA)